

COSC 202 Project

November 28, 2024

1 Data Gathering

We will start by getting the data from the 'fashion_dataset' directory.

```
[ ]: !wget https://github.com/mohamed-12-4/DSAI--COSC-202-/raw/refs/heads/master/  
      ↪fashion_dataset.zip  
      !unzip fashion_dataset
```

```
[19]: import os  
import glob  
import pandas as pd  
import cv2  
import random  
import tensorflow as tf  
import numpy as np  
import matplotlib.pyplot as plt
```

We start by defining a function that takes the path of folder and store the image path and label into a pandas dataframe

```
[20]: def load_data(path: str) -> pd.DataFrame:  
      """  
      Load the data from the path and return a pandas dataframe with the images_  
      ↪path and labels.  
      param path: str: path to the images directory  
      """  
      images = []  
      labels = []  
      for image_path in glob.glob(path + '/*/*.png'):  
          images.append(image_path)  
          labels.append(image_path.split('/')[-2])  
  
      images_df = pd.DataFrame({'images_path': images, 'labels': labels})  
      return images_df
```

```
[21]: images_df = load_data('./fashion_dataset')
```

2 Data Exploration and Visualization

```
[22]: images_df
```

```
[22]:
```

	images_path	labels
0	./fashion_dataset/3/42968.png	3
1	./fashion_dataset/3/28724.png	3
2	./fashion_dataset/3/13600.png	3
3	./fashion_dataset/3/32667.png	3
4	./fashion_dataset/3/32859.png	3
...
59995	./fashion_dataset/5/49155.png	5
59996	./fashion_dataset/5/3531.png	5
59997	./fashion_dataset/5/6438.png	5
59998	./fashion_dataset/5/58969.png	5
59999	./fashion_dataset/5/6929.png	5

```
[60000 rows x 2 columns]
```

```
[23]: images_df.labels = images_df.labels.astype(int)
      images_df.labels.value_counts()
```

```
[23]: labels
3      6000
4      6000
8      6000
7      6000
9      6000
0      6000
6      6000
2      6000
1      6000
5      6000
Name: count, dtype: int64
```

As we can see we have exactly 6,000 images per class which implies that the dataset is balanced

```
[24]: images_df
```

```
[24]:
```

	images_path	labels
0	./fashion_dataset/3/42968.png	3
1	./fashion_dataset/3/28724.png	3
2	./fashion_dataset/3/13600.png	3
3	./fashion_dataset/3/32667.png	3
4	./fashion_dataset/3/32859.png	3
...
59995	./fashion_dataset/5/49155.png	5

```

59996    ./fashion_dataset/5/3531.png        5
59997    ./fashion_dataset/5/6438.png        5
59998    ./fashion_dataset/5/58969.png       5
59999    ./fashion_dataset/5/6929.png        5

```

```
[60000 rows x 2 columns]
```

By inspection, we can convert each label to a name class for an easier interpretation of that data

```

[25]: labels_map = {
        0: 'T-shirt/top',
        1: 'Trouser',
        2: 'Pullover',
        3: 'Dress',
        4: 'Coat',
        5: 'Sandal',
        6: 'Shirt',
        7: 'Sneaker',
        8: 'Bag',
        9: 'Ankle boot'
    }

    def class_to_name(label) -> str:
        """
        Convert the label to the name of the class
        param label: int: label of the class
        """

        return labels_map[label]

```

We also created a function that views an image given its path

```


[26]: def view_image(path: str):
        """
        View the image from the path
        param path: str: path to the image
        """

        import matplotlib.pyplot as plt
        import matplotlib.image as mpimg
        img = cv2.imread(path)
        plt.imshow(img)

    index = random.randint(0, len(images_df))
    image_path = images_df.images_path[index]
    label = images_df.labels[index]
    print(class_to_name(label))
    view_image(image_path)

```


Sneaker



```
[27]: # plot a pie plot of the classes
images_df['class_name'] = images_df['labels'].apply(class_to_name)

images_df.class_name.value_counts().plot(kind='pie', autopct='%1.1f%%', )
```

```
[27]: <Axes: ylabel='count'>
```



3 Data Preprocessing

We start by utilising `tf.ImageDataGenerator`, then we apply `flow_from_dataframe` function, this will go to each file path in the images dataframe and read it, then it will save it in a dataset, we also have the option to choose the batch size.

```
[28]: def create_dataset(df: pd.DataFrame, batch_size: int=32, test=False) -> tf.data.
Dataset:
"""
Create the dataset from the dataframe
param df: pd.DataFrame: dataframe with the images path and labels
param batch_size: int: batch size of the dataset
"""
imagegen = tf.keras.preprocessing.image.ImageDataGenerator(rescale = 1/255.)
if not test:
    dataset = imagegen.flow_from_dataframe(
        df,
        x_col='images_path',
        y_col='labels',
        target_size=(28, 28),
        class_mode='raw',
        batch_size=batch_size,
    )
```

```

else:
    dataset = imagegen.flow_from_dataframe(
        df,
        x_col='images_path',
        y_col='labels',
        target_size=(28, 28),
        class_mode='raw',
        batch_size=1,
        shuffle=False
    )
return dataset

```

```

[29]: from sklearn.model_selection import train_test_split
train_df, test_df = train_test_split(images_df, test_size=0.2, random_state=42)
test_df, val_df = train_test_split(test_df, test_size=0.5, random_state=42)

```

Then we start by applying the create_dataset function to the train, test, validation splits of the dataframe

```

[30]: dataset = create_dataset(images_df, 64)
train_dataset = create_dataset(train_df, 64)
test_dataset = create_dataset(test_df, 1, test=True)
val_dataset = create_dataset(val_df, 64)

```

```

Found 60000 validated image filenames.
Found 48000 validated image filenames.
Found 6000 validated image filenames.
Found 6000 validated image filenames.

```

```

[31]: len(dataset) # 938 batches each batch contains 64 images

```

```

[31]: 938

```

```

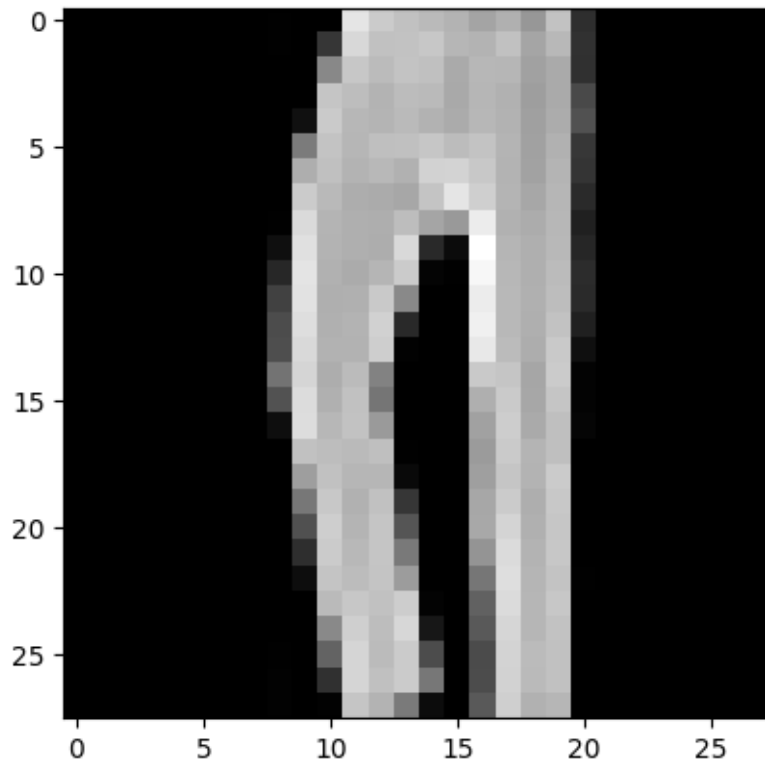
[39]: plt.imshow(dataset[0][0][0])

```

```

[39]: <matplotlib.image.AxesImage at 0x7f024fa19760>

```



4 Simple NN (Only dense layers)

We start by defining a simple model that only contains Dense layers

```
[ ]: Dense_model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28, 3)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
Dense_model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	
Param #		
flatten_3 (Flatten)	(None, 2352)	
↪ 0		

```
dense_8 (Dense)                (None, 128)
↳301,184
dense_9 (Dense)                (None, 64)
↳8,256
dense_10 (Dense)               (None, 10)
↳650
```

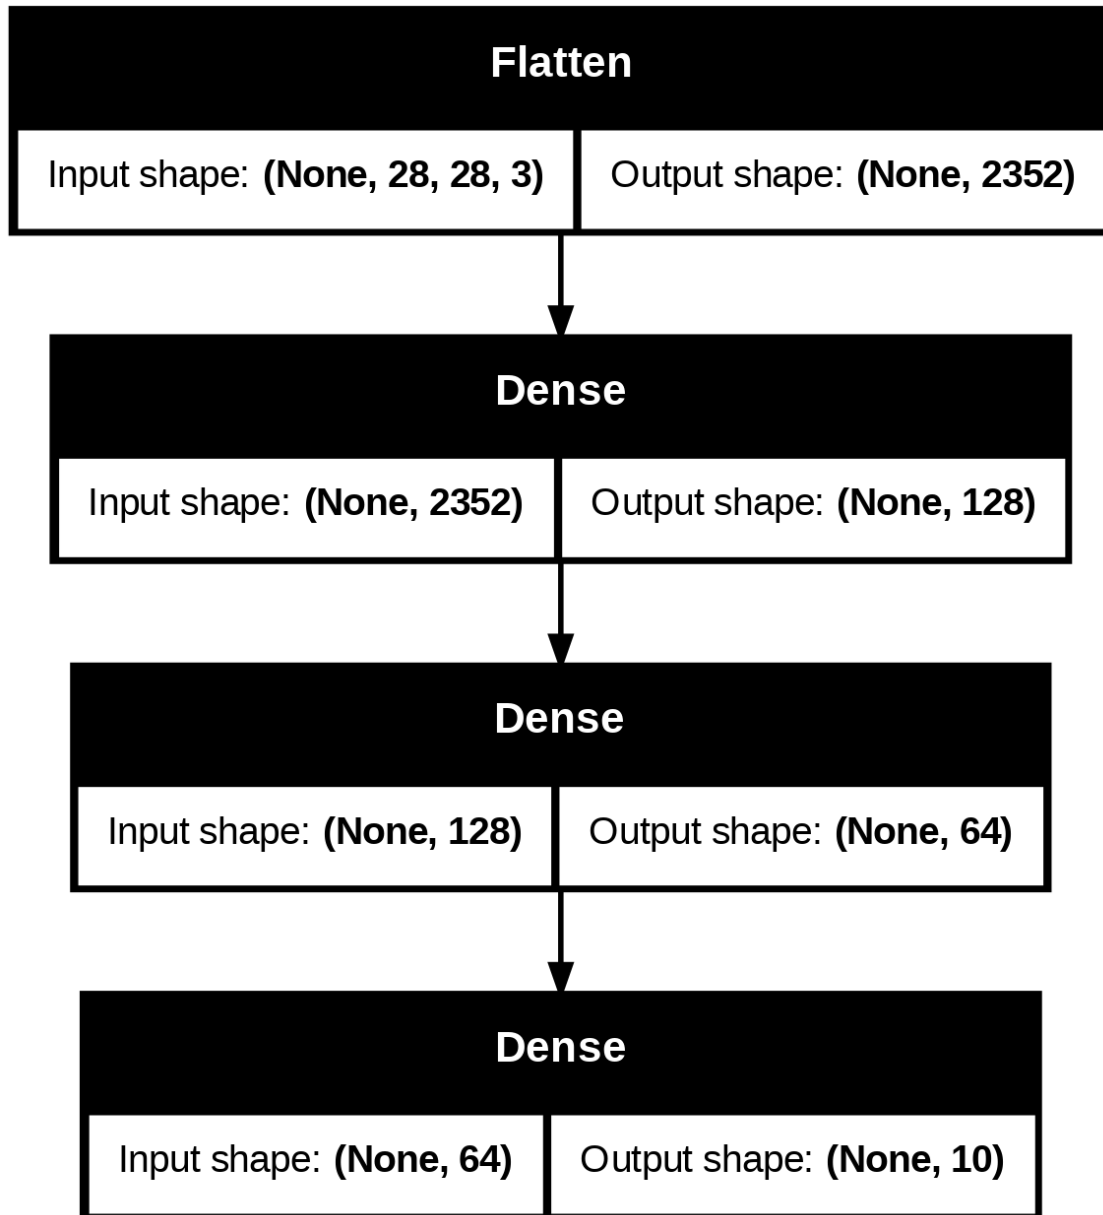
Total params: 310,090 (1.18 MB)

Trainable params: 310,090 (1.18 MB)

Non-trainable params: 0 (0.00 B)

```
[ ]: tf.keras.utils.plot_model(Dense_model, show_shapes=True)
```

```
[ ]:
```

```
[ ]: # adding a callback to prevent the model from overfitting
callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5,
↳restore_best_weights=True)
Dense_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
↳metrics=['accuracy'])
history = Dense_model.fit(train_dataset, epochs=100, callbacks=[callback],
↳validation_data=val_dataset)
```

Epoch 1/100

750/750

19s 23ms/step -

accuracy: 0.7602 - loss: 0.6916 - val_accuracy: 0.8510 - val_loss: 0.4319
 Epoch 2/100
 750/750 22s 26ms/step -
 accuracy: 0.8542 - loss: 0.4024 - val_accuracy: 0.8562 - val_loss: 0.3952
 Epoch 3/100
 750/750 16s 22ms/step -
 accuracy: 0.8692 - loss: 0.3575 - val_accuracy: 0.8658 - val_loss: 0.3747
 Epoch 4/100
 750/750 21s 22ms/step -
 accuracy: 0.8779 - loss: 0.3340 - val_accuracy: 0.8715 - val_loss: 0.3449
 Epoch 5/100
 750/750 17s 23ms/step -
 accuracy: 0.8791 - loss: 0.3202 - val_accuracy: 0.8700 - val_loss: 0.3567
 Epoch 6/100
 750/750 17s 23ms/step -
 accuracy: 0.8851 - loss: 0.3037 - val_accuracy: 0.8738 - val_loss: 0.3505
 Epoch 7/100
 750/750 20s 23ms/step -
 accuracy: 0.8902 - loss: 0.2937 - val_accuracy: 0.8788 - val_loss: 0.3398
 Epoch 8/100
 750/750 20s 22ms/step -
 accuracy: 0.8912 - loss: 0.2859 - val_accuracy: 0.8720 - val_loss: 0.3629
 Epoch 9/100
 750/750 20s 26ms/step -
 accuracy: 0.8980 - loss: 0.2708 - val_accuracy: 0.8672 - val_loss: 0.3667
 Epoch 10/100
 750/750 25s 33ms/step -
 accuracy: 0.8973 - loss: 0.2714 - val_accuracy: 0.8837 - val_loss: 0.3300
 Epoch 11/100
 750/750 25s 33ms/step -
 accuracy: 0.9024 - loss: 0.2560 - val_accuracy: 0.8812 - val_loss: 0.3328
 Epoch 12/100
 750/750 19s 25ms/step -
 accuracy: 0.9077 - loss: 0.2435 - val_accuracy: 0.8847 - val_loss: 0.3293
 Epoch 13/100
 750/750 21s 28ms/step -
 accuracy: 0.9089 - loss: 0.2406 - val_accuracy: 0.8823 - val_loss: 0.3397
 Epoch 14/100
 750/750 39s 25ms/step -
 accuracy: 0.9131 - loss: 0.2321 - val_accuracy: 0.8845 - val_loss: 0.3406
 Epoch 15/100
 750/750 18s 23ms/step -
 accuracy: 0.9113 - loss: 0.2294 - val_accuracy: 0.8840 - val_loss: 0.3404
 Epoch 16/100
 750/750 16s 21ms/step -
 accuracy: 0.9128 - loss: 0.2245 - val_accuracy: 0.8815 - val_loss: 0.3610
 Epoch 17/100
 750/750 17s 23ms/step -

accuracy: 0.9175 - loss: 0.2198 - val_accuracy: 0.8798 - val_loss: 0.3585

As we can see the mode has a validation accuracy of around 88%

5 Convolutional NN

6 Convolutional NN w/MaxPooling layers

Now, we try Convolutional Neural Networks on the data set, we start by a simple one of only Conv and MaxPooling Layers

```
[ ]: model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='leaky_relu',
    ↪input_shape=(28, 28, 3)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='leaky_relu'),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='leaky_relu'),
    tf.keras.layers.Conv2D(64, (3, 3), activation='leaky_relu'),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(64, activation='leaky_relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.summary()
```

```
/usr/local/lib/python3.10/dist-
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential_1"

Layer (type)	Output Shape	
↪Param #		
conv2d (Conv2D)	(None, 26, 26, 32)	↪
↪896		
conv2d_1 (Conv2D)	(None, 24, 24, 64)	↪
↪18,496		
max_pooling2d (MaxPooling2D)	(None, 12, 12, 64)	↪
↪ 0		

conv2d_2 (Conv2D)	(None, 10, 10, 64)	└
↳ 36,928		
conv2d_3 (Conv2D)	(None, 8, 8, 64)	└
↳ 36,928		
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 64)	└
↳ 0		
flatten_1 (Flatten)	(None, 1024)	└
↳ 0		
dense_3 (Dense)	(None, 64)	└
↳ 65,600		
dense_4 (Dense)	(None, 10)	└
↳ 650		

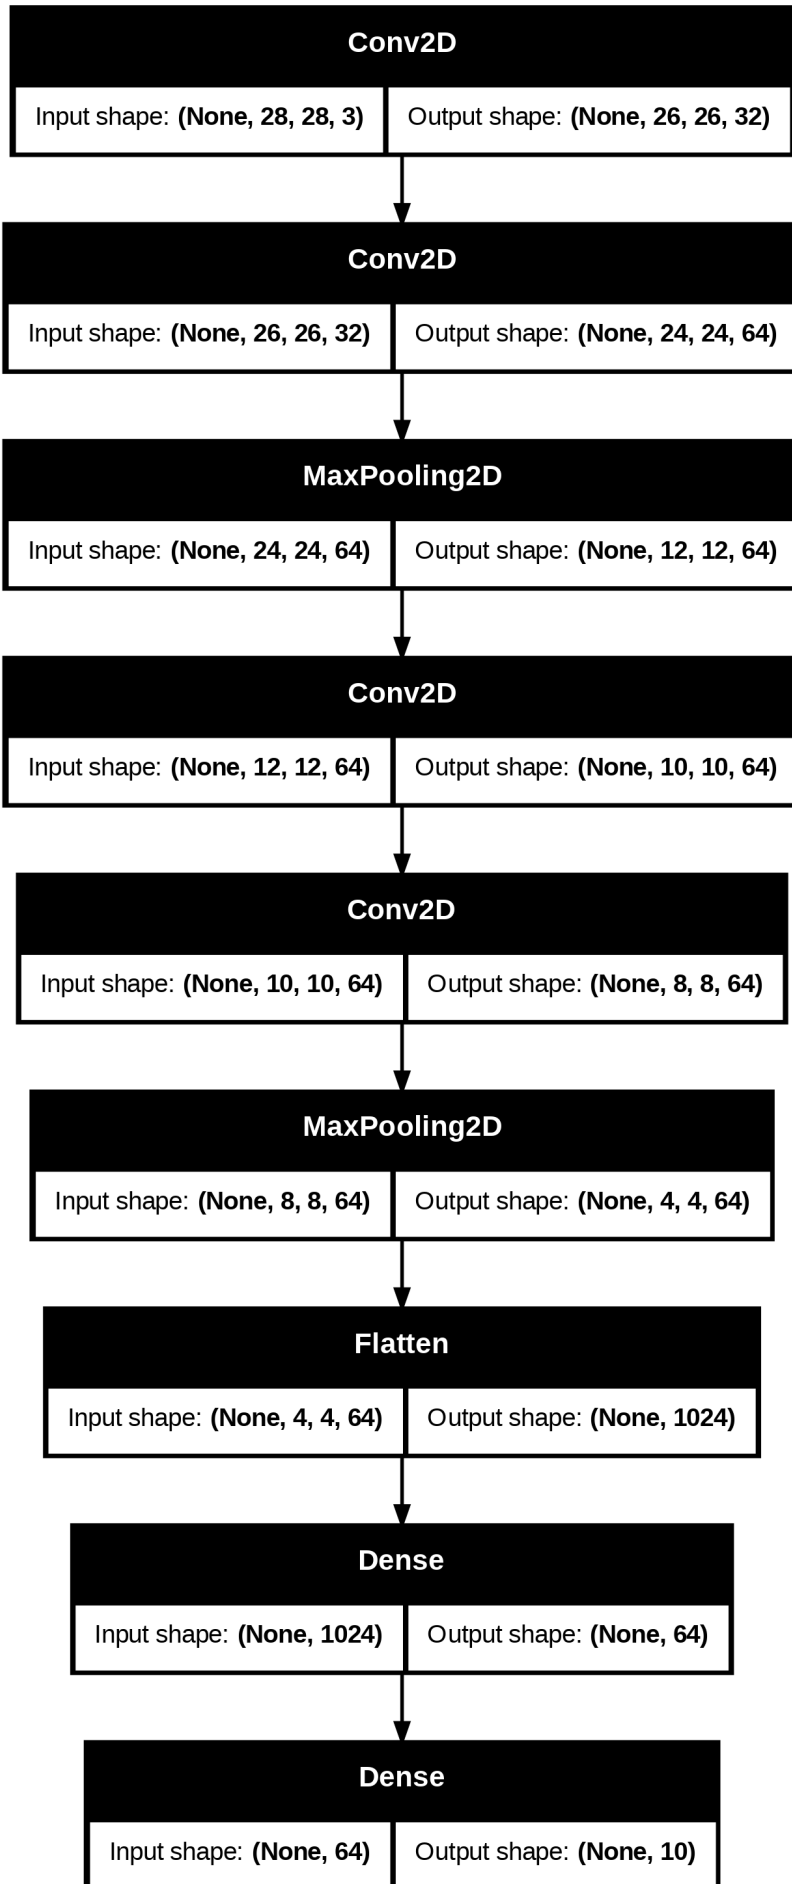
Total params: 159,498 (623.04 KB)

Trainable params: 159,498 (623.04 KB)

Non-trainable params: 0 (0.00 B)

```
[ ]: tf.keras.utils.plot_model(model, show_shapes=True)
```

```
[ ]:
```



```
[ ]: callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5,
↳restore_best_weights=True)
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
↳metrics=['accuracy'])
history = model.fit(train_dataset, epochs=100, callbacks=[callback],
↳validation_data=val_dataset)
```

Epoch 1/100

/usr/local/lib/python3.10/dist-

packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:122:

UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.

self._warn_if_super_not_called()

750/750 25s 26ms/step -

accuracy: 0.7476 - loss: 0.6945 - val_accuracy: 0.8727 - val_loss: 0.3612

Epoch 2/100

750/750 36s 25ms/step -

accuracy: 0.8865 - loss: 0.3103 - val_accuracy: 0.8967 - val_loss: 0.2875

Epoch 3/100

750/750 19s 25ms/step -

accuracy: 0.9057 - loss: 0.2585 - val_accuracy: 0.9005 - val_loss: 0.2771

Epoch 4/100

750/750 19s 23ms/step -

accuracy: 0.9186 - loss: 0.2271 - val_accuracy: 0.9088 - val_loss: 0.2610

Epoch 5/100

750/750 22s 24ms/step -

accuracy: 0.9253 - loss: 0.1995 - val_accuracy: 0.8947 - val_loss: 0.2823

Epoch 6/100

750/750 17s 23ms/step -

accuracy: 0.9333 - loss: 0.1808 - val_accuracy: 0.9020 - val_loss: 0.2874

Epoch 7/100

750/750 21s 24ms/step -

accuracy: 0.9399 - loss: 0.1653 - val_accuracy: 0.9125 - val_loss: 0.2562

Epoch 8/100

750/750 17s 23ms/step -

accuracy: 0.9434 - loss: 0.1496 - val_accuracy: 0.9165 - val_loss: 0.2507

Epoch 9/100

750/750 21s 24ms/step -

accuracy: 0.9531 - loss: 0.1284 - val_accuracy: 0.9098 - val_loss: 0.2877

Epoch 10/100

750/750 17s 23ms/step -

accuracy: 0.9567 - loss: 0.1166 - val_accuracy: 0.9082 - val_loss: 0.2751

Epoch 11/100

750/750 18s 24ms/step -

accuracy: 0.9600 - loss: 0.1045 - val_accuracy: 0.9077 - val_loss: 0.3118

Epoch 12/100

750/750 18s 24ms/step -

accuracy: 0.9652 - loss: 0.0935 - val_accuracy: 0.9095 - val_loss: 0.3285


Epoch 13/100

750/750 20s 23ms/step -

accuracy: 0.9697 - loss: 0.0831 - val_accuracy: 0.9050 - val_loss: 0.3681

We noticed that this model perform slightly better on the data acheiving around 90% in validation accuracy

```
[ ]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```



7 Convolutional NN + Regulization / Normalization layers

Now we Try a better one that have an extra Regulization (Dropout layers) and Normalization Layers. These layers will prevent the model from overfitting.

```
[ ]: CNNR_model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='leaky_relu',
    ↪input_shape=(28, 28, 3)),
    tf.keras.layers.Conv2D(32, (3, 3), activation='leaky_relu'),
    tf.keras.layers.Normalization(),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='leaky_relu'),
    tf.keras.layers.Conv2D(64, (3, 3), activation='leaky_relu'),
    tf.keras.layers.Normalization(),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(64, activation='leaky_relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
CNNR_model.summary()
```

```
/usr/local/lib/python3.10/dist-
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
Model: "sequential_7"
```

Layer (type)	Output Shape	
↪Param #		
conv2d_43 (Conv2D)	(None, 26, 26, 32)	↪
↪896		
conv2d_44 (Conv2D)	(None, 24, 24, 32)	↪
↪9,248		
normalization (Normalization)	(None, 24, 24, 32)	↪
↪ 65		
dropout (Dropout)	(None, 24, 24, 32)	↪
↪ 0		

max_pooling2d_17 (MaxPooling2D)	(None, 12, 12, 32)	└
↪ 0		
conv2d_45 (Conv2D)	(None, 10, 10, 64)	└
↪ 18,496		
conv2d_46 (Conv2D)	(None, 8, 8, 64)	└
↪ 36,928		
normalization_1 (Normalization)	(None, 8, 8, 64)	└
↪ 129		
dropout_1 (Dropout)	(None, 8, 8, 64)	└
↪ 0		
max_pooling2d_18 (MaxPooling2D)	(None, 4, 4, 64)	└
↪ 0		
flatten_7 (Flatten)	(None, 1024)	└
↪ 0		
dense_20 (Dense)	(None, 64)	└
↪ 65,600		
dense_21 (Dense)	(None, 10)	└
↪ 650		

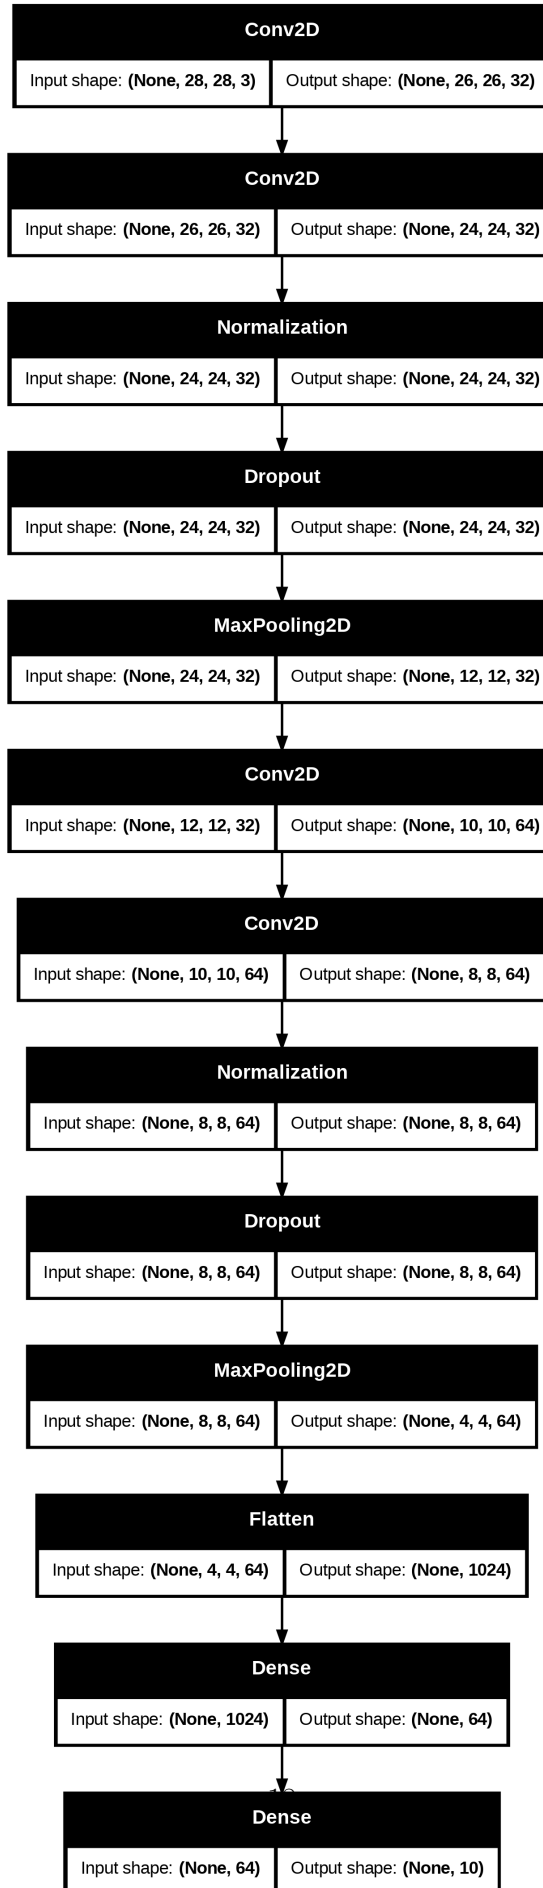
Total params: 132,012 (515.68 KB)

Trainable params: 131,818 (514.91 KB)

Non-trainable params: 194 (784.00 B)

```
[ ]: tf.keras.utils.plot_model(CNNR_model, show_shapes=True)
```

```
[ ]:
```



```
[ ]: callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=5,
↳restore_best_weights=True)
CNNR_model.compile(optimizer='nadam', loss='sparse_categorical_crossentropy',
↳metrics=['accuracy'])
history = CNNR_model.fit(train_dataset, epochs=100, callbacks=[callback],
↳validation_data=val_dataset)
```

```
Epoch 1/100
750/750          24s 26ms/step -
accuracy: 0.7501 - loss: 0.6978 - val_accuracy: 0.8693 - val_loss: 0.3765
Epoch 2/100
750/750          37s 25ms/step -
accuracy: 0.8802 - loss: 0.3223 - val_accuracy: 0.8880 - val_loss: 0.3129
Epoch 3/100
750/750          20s 25ms/step -
accuracy: 0.8980 - loss: 0.2743 - val_accuracy: 0.9002 - val_loss: 0.2799
Epoch 4/100
750/750          19s 25ms/step -
accuracy: 0.9084 - loss: 0.2476 - val_accuracy: 0.9002 - val_loss: 0.2788
Epoch 5/100
750/750          18s 23ms/step -
accuracy: 0.9145 - loss: 0.2320 - val_accuracy: 0.9097 - val_loss: 0.2520
Epoch 6/100
750/750          18s 23ms/step -
accuracy: 0.9229 - loss: 0.2164 - val_accuracy: 0.9075 - val_loss: 0.2609
Epoch 7/100
750/750          21s 23ms/step -
accuracy: 0.9238 - loss: 0.2027 - val_accuracy: 0.9122 - val_loss: 0.2487
Epoch 8/100
750/750          22s 25ms/step -
accuracy: 0.9296 - loss: 0.1903 - val_accuracy: 0.9125 - val_loss: 0.2508
Epoch 9/100
750/750          17s 23ms/step -
accuracy: 0.9353 - loss: 0.1751 - val_accuracy: 0.9147 - val_loss: 0.2432
Epoch 10/100
750/750          21s 24ms/step -
accuracy: 0.9365 - loss: 0.1675 - val_accuracy: 0.9092 - val_loss: 0.2655
Epoch 11/100
750/750          18s 24ms/step -
accuracy: 0.9415 - loss: 0.1561 - val_accuracy: 0.9158 - val_loss: 0.2395
Epoch 12/100
750/750          17s 23ms/step -
accuracy: 0.9435 - loss: 0.1505 - val_accuracy: 0.9153 - val_loss: 0.2566
Epoch 13/100
750/750          18s 24ms/step -
```

```
accuracy: 0.9481 - loss: 0.1405 - val_accuracy: 0.9137 - val_loss: 0.2683
Epoch 14/100
```

```
750/750          17s 23ms/step -
```

```
accuracy: 0.9491 - loss: 0.1362 - val_accuracy: 0.9092 - val_loss: 0.2682
Epoch 15/100
```

```
750/750          17s 23ms/step -
```

```
accuracy: 0.9518 - loss: 0.1281 - val_accuracy: 0.9138 - val_loss: 0.2662
Epoch 16/100
```

```
750/750          18s 24ms/step -
```

```
accuracy: 0.9503 - loss: 0.1323 - val_accuracy: 0.9155 - val_loss: 0.2730
```

We can see that this model is better than the previous two with 91.5% validation accuracy

```
[ ]: CNNR_model.evaluate(test_dataset)
```

```
94/94           2s 18ms/step -
```

```
accuracy: 0.9234 - loss: 0.2260
```

```
[ ]: [0.23158198595046997, 0.9211666584014893]
```

We will choose the last model as our best one

Lastly we save the model to be tested on real-life data on another notebook file

```
[ ]: import pickle
with open('model.pkl', 'wb') as f:
    pickle.dump(CNNR_model, f)
```

8 Test

```
[10]: import pickle
model = pickle.load(open('model.pkl', 'rb'))
```

```
W0000 00:00:1732646876.483687 37084 gpu_device.cc:2344] Cannot dlopen some GPU
libraries. Please make sure the missing libraries mentioned above are installed
properly if you would like to use GPU. Follow the guide at
https://www.tensorflow.org/install/gpu for how to download and setup the
required libraries for your platform.
Skipping registering GPU devices...
```

```
[2]: model.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
conv2d_43 (Conv2D)	(None, 26, 26, 32)	896

conv2d_44 (Conv2D)	(None, 24, 24, 32)	9,248
normalization (Normalization)	(None, 24, 24, 32)	65
dropout (Dropout)	(None, 24, 24, 32)	0
max_pooling2d_17 (MaxPooling2D)	(None, 12, 12, 32)	0
conv2d_45 (Conv2D)	(None, 10, 10, 64)	18,496
conv2d_46 (Conv2D)	(None, 8, 8, 64)	36,928
normalization_1 (Normalization)	(None, 8, 8, 64)	129
dropout_1 (Dropout)	(None, 8, 8, 64)	0
max_pooling2d_18 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten_7 (Flatten)	(None, 1024)	0
dense_20 (Dense)	(None, 64)	65,600
dense_21 (Dense)	(None, 10)	650

Total params: 395,651 (1.51 MB)

Trainable params: 131,818 (514.91 KB)

Non-trainable params: 194 (784.00 B)

Optimizer params: 263,639 (1.01 MB)

```
[ ]: import numpy as np
def model_predict(model, data):
    """
    Predict the data using the model
    param model: the model to predict the data
    param data: the data to be used for prediction
    """
    predictions = model.predict(data)
    return np.argmax(predictions, axis=1)
```

```
def compare_predictions(predictions, labels):
    """
    Compare the predictions with the labels and print the accuracy
    param predictions: the predictions of the model
    param labels: the labels of the data
    """
    for i in predictions:
        print('Prediction:', class_to_name(i), end=' ')
        print('Label:', class_to_name(labels[i]))

    print('Accuracy:', np.mean(predictions == labels))
```

```
[42]: model.evaluate(test_dataset)
```

```
94/94          2s 17ms/step -
accuracy: 0.9471 - loss: 0.1382
```

```
[42]: [0.1509178876876831, 0.9453333616256714]
```

Now Applying the model on test data we get 94% accuracy

8.1 Confusion matrix

```
[ ]: # create confusion matrix
from sklearn.metrics import confusion_matrix
import seaborn as sns

y_pred = model_predict(model, test_dataset)
y_true = test_df.labels.values
#
sns.heatmap(confusion_matrix(y_true, y_pred), annot=True, fmt='d')
```

```
6000/6000          12s 2ms/step
```

```
[ ]: <Axes: >
```

